

Review Radar

Leveraging NLP for Targeted Product Enhancement

CS 816- Software Production Engineering



International Institute of Information Technology, Bangalore

under the guidance of

Prof. B Thangaraju

Submitted by

Anjali J Kumar (MT2023126)

Sambhu S S (MT2023072)

[GitHub](#)

CONTENTS

Table of Contents

ABSTRACT	3
PROJECT DETAILS	4
DevOps	8
DEVELOPMENT ENVIRONMENT	12
TEST	14
SOURCE CODE MANAGEMENT	15
CONTINUOUS INTEGRATION WITH JENKINS	17
CONTAINERIZATION WITH DOCKER	31
DEPLOYMENT WITH KUBERNETES	34
CONFIGURATION MANAGEMENT WITH ANSIBLE	40
CONTINUOUS MONITORING WITH THE ELK STACK	44
CONCLUSION	47

ABSTRACT

In the realm of consumer electronics, understanding customer sentiments and identifying key areas for improvement is paramount for staying competitive. This project presents a comprehensive solution leveraging DevOps methodologies to enhance the efficiency and reliability of a sentiment analysis and topic modelling system for phone companies.

The project begins with the development phase, wherein a dataset comprising reviews of various phone companies is subjected to topic modelling using Non-Negative Matrix Factorization (NMF). This process identifies latent topics within the reviews, offering insights into the diverse aspects of customer feedback.

Building upon the topic modelling results, sentiment analysis is performed on each review within the identified topics. By gauging the sentiment of individual reviews, the system identifies the areas of concern and excellence for each phone company, facilitating targeted improvements and highlighting strengths.

The core of the project lies in its DevOps implementation, aimed at automating and streamlining the deployment and maintenance processes. Continuous Integration (CI) and Continuous Deployment (CD) pipelines ensure seamless updates and enhancements to the sentiment analysis and topic modelling algorithms. Additionally, robust monitoring and logging mechanisms are integrated to provide real-time insights into system performance and user interactions.

A user-friendly dashboard is provided to phone companies, presenting actionable insights derived from the sentiment analysis and topic modelling results. Moreover, a feedback mechanism allows for the submission of new reviews, which are promptly analysed to provide instant feedback on sentiment and topic relevance.

Through this project, the fusion of advanced data analytics techniques with DevOps practices demonstrates a holistic approach towards enhancing customer experience and driving business success in the dynamic consumer electronics landscape.

PROJECT DETAILS

Introduction

This project aims to analyze customer reviews of Apple and Samsung phones to provide insights into customer sentiments and highlight key areas for improvement. By using topic modelling and sentiment analysis, the system helps companies understand customer feedback comprehensively and take data-driven actions.

Data Collection and Preparation

The initial dataset comprised approximately 400,000 reviews. From this, reviews related to Apple and Samsung were extracted. The preprocessing steps involved:

Cleaning the Text Data: Removing stop words, punctuation, and performing tokenization.

Filtering: Only retaining reviews related to Apple and Samsung for focused analysis.

Topic Modelling

Topic modelling was performed using Non-Negative Matrix Factorization (NMF) to identify the main topics within the reviews. Ten topics were identified, each characterized by the top 15 words associated with them. These topics were manually named for clarity:

- Appreciation Towards Service and Delivery
- Buying Experience
- Condition and Packaging of Items
- Expectations and Product Quality
- Issues with Device Functionality
- Ordering and Customer Service
- Overall Satisfaction with Purchase
- Product Quality and Experience
- Recommendations and Satisfaction
- Satisfaction with Product Condition

Implementation:

NMF Configuration: NMF was configured to generate 10 topics.

Evaluation: Topics were evaluated and named based on the top 15 words identified by the model.

Sentiment Analysis

Sentiment analysis was conducted using the VADER sentiment analysis tool. Each review was classified into one of three categories: positive, neutral, or negative. This analysis provided insights into the overall sentiment distribution across different topics.

Implementation:

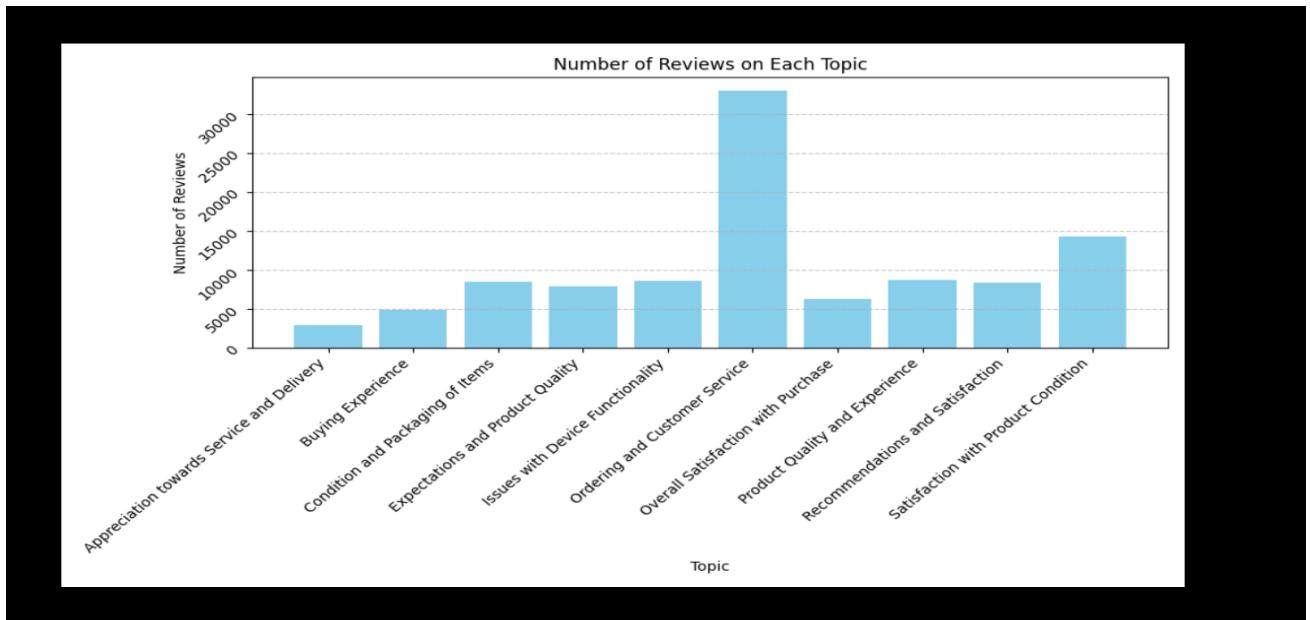
Sentiment Classification: VADER was used to analyze and classify sentiments for each review.

Result Aggregation: Sentiments were aggregated for each topic to understand the distribution.

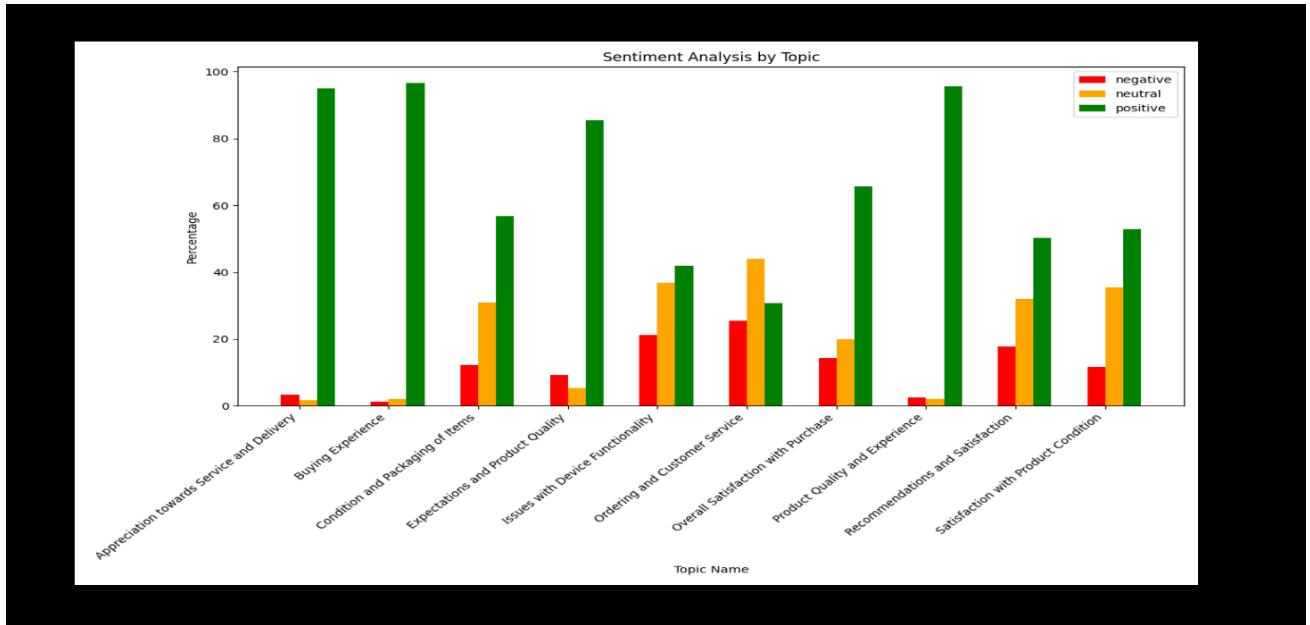
Dashboard

The main features of the dashboard include:

Topic Distribution Bar Plot: This bar plot shows the number of reviews associated with each topic. It helps companies understand which aspects customers frequently comment on.



Sentiment Analysis by Topic: This bar plot displays the percentage of positive, neutral, and negative reviews for each topic. It helps companies identify areas needing attention.



The dashboard provides a clear visual representation of the topics and their corresponding sentiments, making it easy for companies to interpret the data and derive actionable insights.

The firm should focus more on the 'Ordering and Customer Service' aspect, as it has the highest percentage of negative reviews (25.43%). There is scope for improvement in the 'Ordering and Customer Service' aspect, as it has the highest percentage of neutral reviews (43.87%). The firm is performing well in the 'Buying Experience' aspect, as it has the highest percentage of positive reviews (96.58%).

User Interaction

The dashboard includes a text box for users to submit new reviews. The system analyzes these new reviews in real-time, predicting the relevant topic and sentiment. The results are displayed immediately to the user, providing instant feedback.

Add Your Voice

Have something to say about a brand? Share your thoughts by typing a new review into our text area. Review Radar will instantly analyze it, providing insights into the topic and sentiment of your review.

this is a new review on a good phone

Predicted Topic: Expectations and Product Quality

Sentiment: positive

Implementation:

Real-Time Analysis: New reviews are analysed in real-time for topic prediction and sentiment analysis.

Feedback Display: The predicted topic and sentiment are displayed instantly to the user.

Data Storage: Newly submitted reviews are stored in a MySQL database for future analysis, although they are not currently used in the live system.

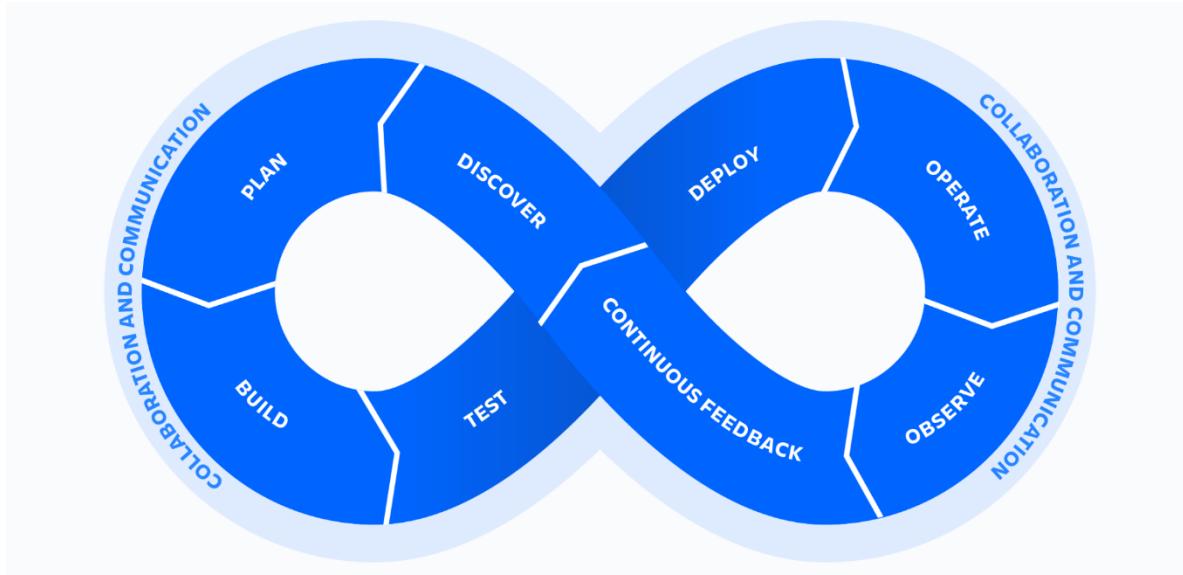
Conclusion

This project effectively utilizes topic modelling and sentiment analysis to provide actionable insights from customer reviews of Apple and Samsung phones. By identifying key topics and their associated sentiments, companies can focus on areas needing improvement and recognize their strengths. The interactive dashboard enhances user experience by allowing real-time analysis of new reviews, ensuring continuous feedback and improvement.

DevOps

DevOps is a set of practices aimed at breaking down silos between software development (Dev) and IT operations (Ops) teams, fostering collaboration, and automating processes throughout the software development lifecycle.

The primary goal of DevOps is to increase the speed, efficiency, and quality of software delivery.



1. Assessment and Planning

- Evaluate current development and operations processes, identifying pain points, bottlenecks, and areas for improvement.

2. Culture and Organization

- Encourage communication and transparency between teams to break down silos and build trust.

3. Automation

- Identify manual and repetitive tasks throughout the software development lifecycle, such as building, testing, deployment, and infrastructure provisioning.

4. Continuous Integration (CI)

- Establish a CI pipeline that automatically builds and tests code changes whenever they are committed to version control.

5. Continuous Delivery or Deployment (CD)

- Extend CI to include continuous delivery practices, enabling automated deployment of code changes to production or staging environments.

6. Monitoring and Feedback

- Implement monitoring and logging solutions to track application performance, infrastructure metrics, and user behaviour in real-time.

Why DevOps?

1. **Faster Time to Market:** By automating processes and breaking down barriers between teams, DevOps enables faster delivery of features and updates to customers, allowing businesses to respond to market demands more quickly.
2. **Improved Quality:** Continuous integration, automated testing, and continuous delivery practices help catch bugs early in the development process, leading to higher-quality software with fewer defects.
3. **Increased Collaboration:** DevOps encourages collaboration between development, operations, and other stakeholders, fostering a culture of shared responsibility and accountability.
4. **Better Scalability and Reliability:** Automation and infrastructure-as-code practices enable teams to easily scale infrastructure and deploy applications across different environments, leading to improved scalability and reliability.
5. **Reduced Risk:** By automating processes and deploying smaller, incremental changes, DevOps reduces the risk of introducing errors or failures into production environments.

Tools used

1. **PyCharm/VSCode** - Project IDEs for backend (PyCharm) and frontend (VSCode) development
2. **No formal build tool (Flask-based project)** - Build Tool
3. **unittest** - Testing Tool
4. **Git** - Source Control Management Tool
5. **Jenkins** - Continuous Integration Tool
6. **Docker** - Containerization Tool
7. **Kubernetes** - Container Orchestration Tool
8. **Ansible** - Configuration Management Tool
9. **ELK Stack** - Continuous Monitoring

Software Development Life Cycle (SDLC)

Software development life cycle (SDLC) is a structured process that is used to design, develop, and test good-quality software. SDLC, or software development life cycle, is a methodology that defines the entire procedure of software development step-by-step.

The goal of the SDLC life cycle model is to deliver high-quality, maintainable software that meets the user's requirements.

SDLC in software engineering models outlines the plan for each stage so that each stage of the software development model can perform its task efficiently to deliver the software at a low cost within a given time frame that meets users' requirements.



Steps Followed in the Project

Set Up Development Environment

Install dependencies using pip and create a virtual environment with virtualenv or venv.

Use PyCharm for backend development and VSCode for frontend development.

Test

Write and run unit tests using **unittest** to ensure code quality and functionality.

Source Code Management

Use Git for version control and managing source code repositories.

Continuous Integration

Set up Jenkins to automate testing and integration processes.

Containerization

Use Docker to create container images for the Flask application, ensuring consistency across different environments.

Configuration Management

Utilize Ansible to manage and automate configuration across various environments.

Deployment

Deploy the containerized application using Kubernetes for orchestration and managing scalability.

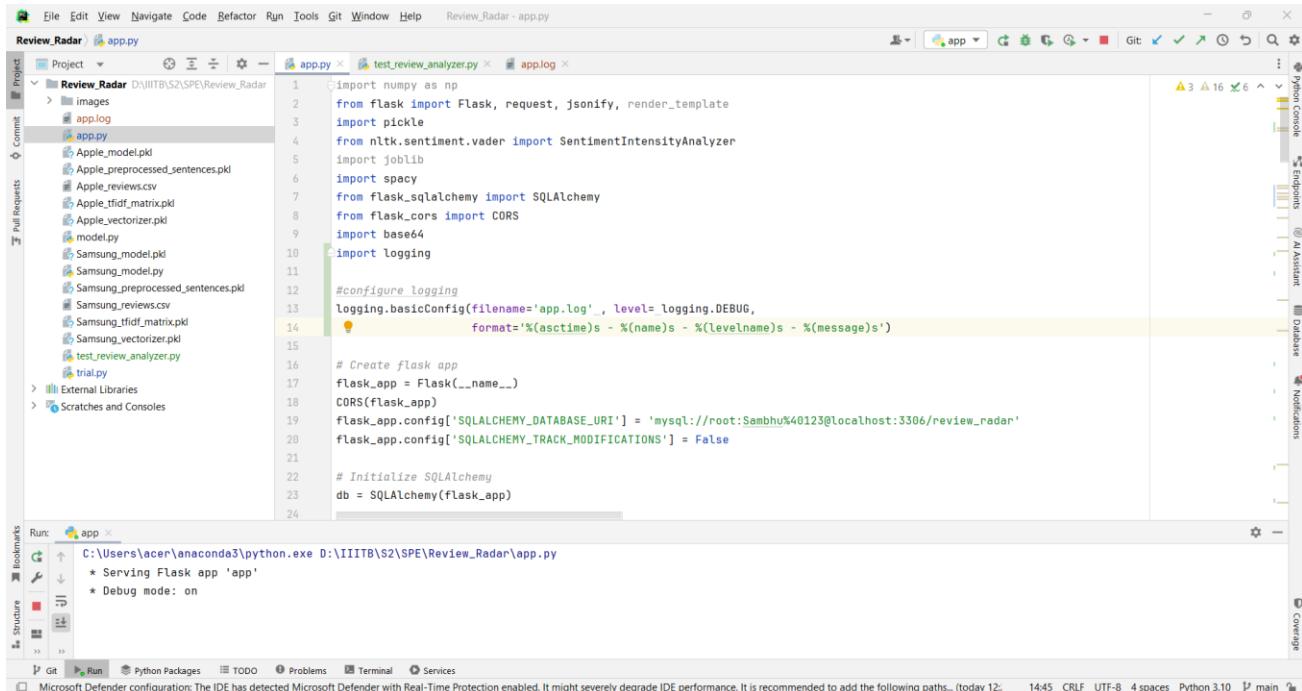
Continuous Monitoring

Implement the ELK Stack (Elasticsearch, Logstash, and Kibana) to monitor application performance and logs continuously.

DEVELOPMENT ENVIRONMENT

Backend Development in PyCharm (IDE):

PyCharm is a feature-rich Integrated Development Environment (IDE) for Python development. It offers tools for code completion, debugging, and version control integration, streamlining the development workflow.



The screenshot shows the PyCharm IDE interface. The project structure on the left includes a 'Review_Radar' folder containing 'app.py', 'app.log', and several preprocessed CSV and PKL files for Apple and Samsung reviews. The code editor on the right displays the 'app.py' file, which imports various libraries like numpy, flask, pickle, nltk, joblib, spaCy, SQLAlchemy, and CORS. It configures logging, creates a Flask app, and initializes SQLAlchemy. The bottom status bar shows the command line running 'python app.py' and the application is serving at 'http://127.0.0.1:5000/'. The bottom right corner shows the Python version as 3.10.

```

import numpy as np
from flask import Flask, request, jsonify, render_template
import pickle
from nltk.sentiment.vader import SentimentIntensityAnalyzer
import joblib
import spaCy
from flask_sqlalchemy import SQLAlchemy
from flask_cors import CORS
import base64
import logging

#configure logging
logging.basicConfig(filename='app.log', level= logging.DEBUG,
                    format='%(asctime)s - %(name)s - %(levelname)s - %(message)s')

# Create flask app
flask_app = Flask(__name__)
CORS(flask_app)
flask_app.config['SQLALCHEMY_DATABASE_URI'] = 'mysql://root:Sambhu%40123@localhost:3306/review_radar'
flask_app.config['SQLALCHEMY_TRACK_MODIFICATIONS'] = False

# Initialize SQLAlchemy
db = SQLAlchemy(flask_app)

```

Dependencies Installed with pip:

Project dependencies were installed using pip, ensuring access to required libraries for the Flask application.

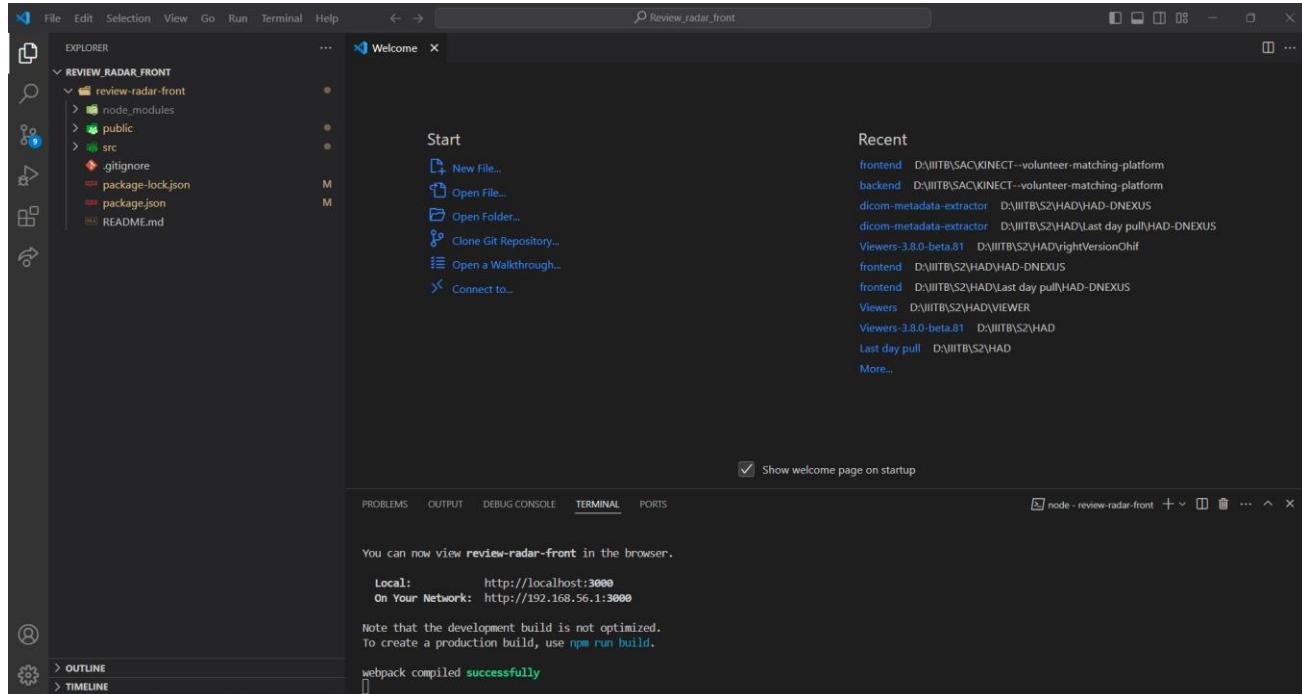
Libraries Used:

- Flask
- NLTK (Natural Language Toolkit) - Specifically, the SentimentIntensityAnalyzer module
- Joblib
- spaCy
- SQLAlchemy
- Flask-CORS
- Base64
- Logging

These libraries serve various purposes in the project, such as web development (Flask), sentiment analysis (NLTK), database interaction (SQLAlchemy), etc.

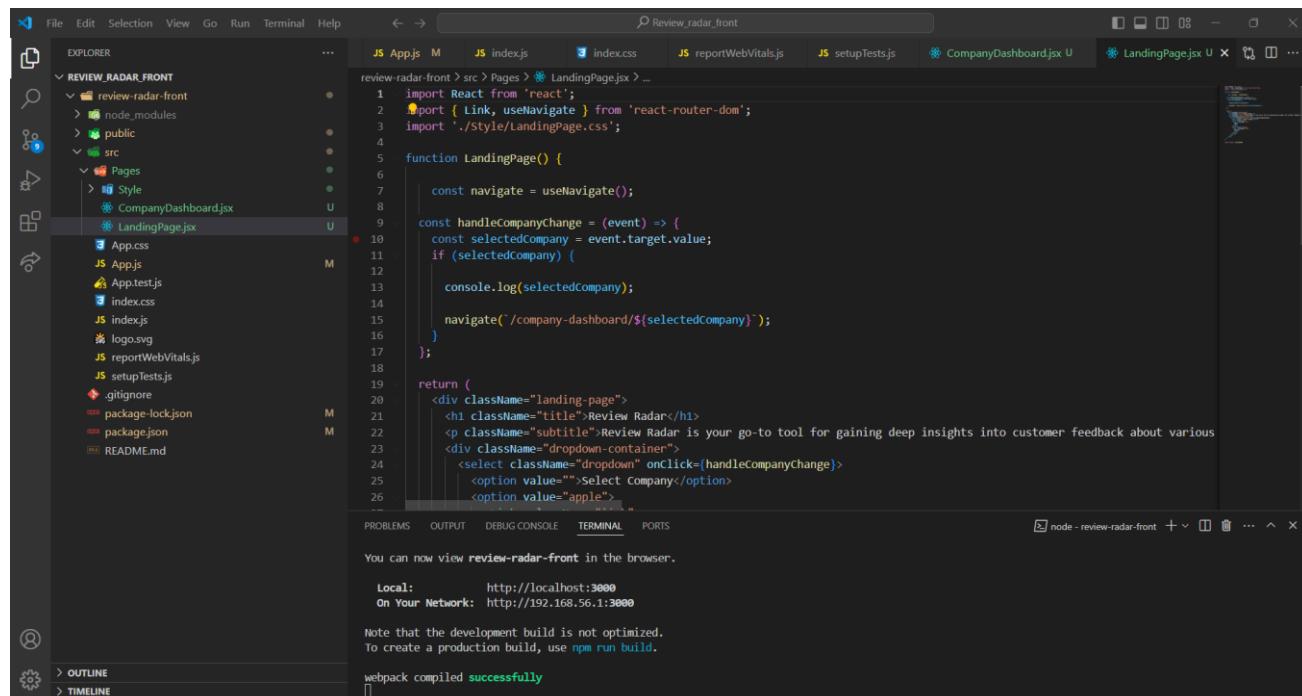
Frontend Development in Visual Studio Code (VSCode):

Visual Studio Code (VSCode) is a versatile code editor used for frontend development, offering features like IntelliSense, debugging, and Git integration.



React:

The frontend of the application was developed using React, a JavaScript library for building user interfaces.



TEST

Test Cases Using unittest Module:

Testing of the Flask application is performed using the unittest module, which provides a framework for organizing and running test cases.

The screenshot shows the PyCharm IDE interface with the following details:

- Project:** Review_Radar
- File:** test_review_analyzer.py
- Code Editor:** The code for `test_review_analyzer.py` is displayed, containing unit tests for the `TestReviewAnalyzer` class.
- Run Tab:** Shows the output of the test run, indicating 2 tests passed in 0.233 ms.
- Bottom Status Bar:** Displays "Tests passed: 2 (moments ago)" and the Python version "Python 3.10".

- The **TestReviewAnalyzer** class contains test cases for analysing reviews of both Apple and Samsung products.
 - The **setUp** and **tearDown** methods ensure that the Flask application context is set up and torn down before and after each test, respectively.
 - Within each test method, reviews are simulated for both Apple and Samsung products, and the **analyze_review** function is called with the simulated data.
 - Assertions are made to verify that the predicted topic and sentiment match the expected values for each review.

SOURCE CODE MANAGEMENT

Git

Git is a distributed version control system used for tracking changes in source code during software development. It allows multiple developers to collaborate on projects efficiently by providing features like branching, merging, and version history tracking.

Repository Setup:

Begin by initializing a Git repository in your project directory using the git init command. This creates a .git directory to store repository metadata.

Committing Changes:

After making changes to your code, use the git add command to stage the modified files for commit. Then, commit the changes to the repository using git commit -m "commit message".

```
anjali@anjali-HP-Pavilion-Laptop-14-ec1xxx:~/SPE_FINAL_PROJECT/REVIEW-RADAR$ git branch
* main
anjali@anjali-HP-Pavilion-Laptop-14-ec1xxx:~/SPE_FINAL_PROJECT/REVIEW-RADAR$ git status
On branch main
Your branch is up to date with 'origin/main'.

Changes not staged for commit:
  (use "git add <file>..." to update what will be committed)
  (use "git restore <file>..." to discard changes in working directory)
    modified:   Jenkinsfile

Untracked files:
  (use "git add <file>..." to include in what will be committed)
    backend/.idea/
    backend/Review_Radar/.idea/
    backend/Review_Radar/app.log
    backend/Review_Radar/app.log.1
    backend/Review_Radar/app.log.10
    backend/Review_Radar/app.log.2
    backend/Review_Radar/app.log.3
    backend/Review_Radar/app.log.4
    backend/Review_Radar/app.log.5
    backend/Review_Radar/app.log.6
    backend/Review_Radar/app.log.7
    backend/Review_Radar/app.log.8
    backend/Review_Radar/app.log.9

no changes added to commit (use "git add" and/or "git commit -a")
anjali@anjali-HP-Pavilion-Laptop-14-ec1xxx:~/SPE_FINAL_PROJECT/REVIEW-RADAR$ git add Jenkinsfile
gitadd: command not found
anjali@anjali-HP-Pavilion-Laptop-14-ec1xxx:~/SPE_FINAL_PROJECT/REVIEW-RADAR$ git add Jenkinsfile
anjali@anjali-HP-Pavilion-Laptop-14-ec1xxx:~/SPE_FINAL_PROJECT/REVIEW-RADAR$ git commit -m "testing webhook"
[main e6f93f5] testing webhook
 1 file changed, 1 insertion(+), 1 deletion(-)
anjali@anjali-HP-Pavilion-Laptop-14-ec1xxx:~/SPE_FINAL_PROJECT/REVIEW-RADAR$ git push -u origin main
Username for 'https://github.com': kumaranjalij
Password for 'https://kumaranjalij@github.com':
Enumerating objects: 5, done.
Counting objects: 100% (5/5), done.
Delta compression using up to 12 threads
Compressing objects: 100% (3/3), done.
Writing objects: 100% (3/3), 329 bytes | 329.00 KiB/s, done.
Total 3 (delta 2), reused 0 (delta 0), pack-reused 0
remote: Resolving deltas: 100% (2/2), completed with 2 local objects.
To https://github.com/kumaranjalij/ReviewRadar.git
  34b4fc..e6f93f5  main -> main
branch 'main' set up to track 'origin/main'.
```

Version History and Log:

Git maintains a detailed version history of your project, allowing you to view commit logs, changes, and contributors. Use commands like git log to view commit history and git diff to compare changes between commits.

GitHub repository screenshot:

The screenshot shows a GitHub repository page for 'ReviewRadar'. The repository is public and has 1 branch and 0 tags. The main branch is 'main'. There are 52 commits listed, with the most recent being a file named 'changing log file size' by user 'kumaranjalij' made 1 hour ago. Other commits include 'ansible-deploy', 'backend/Review_Radar', 'frontend/reviewradar', 'k8s', 'Jenkinsfile', and 'README.md'. The 'Code' tab is selected. On the right side, there is an 'About' section with a note: 'No description, website, or topics provided.' It also lists 'Readme', 'Activity', '0 stars', '1 watching', '0 forks', and 'Report repository'. Below that is a 'Releases' section with 'No releases published'. Under 'Packages', it says 'No packages published'. At the bottom, there is a 'Languages' section.

CONTINUOUS INTEGRATION WITH JENKINS

Jenkins is an open-source automation server that facilitates continuous integration and continuous delivery (CI/CD) pipelines. It automates the build, test, and deployment processes, allowing teams to integrate code changes frequently and detect issues early in the development cycle. Jenkins provides a highly extensible architecture with a vast ecosystem of plugins, enabling customization and integration with various tools and technologies.

Key Features of Jenkins:

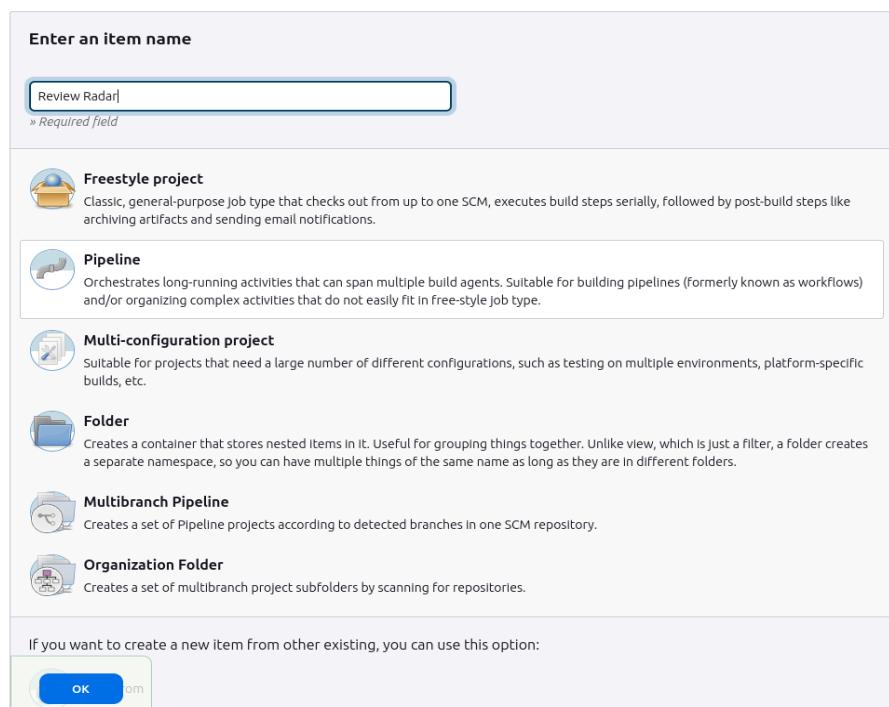
- **Automation:** Jenkins automates repetitive tasks such as code compilation, testing, and deployment, reducing manual effort and improving efficiency.
- **Scalability:** Jenkins scales easily to accommodate projects of any size, from small teams to large enterprises, supporting parallel and distributed builds.
- **Integration:** Jenkins integrates seamlessly with version control systems, issue trackers, and other development tools, enabling end-to-end automation of the software delivery pipeline.
- **Extensibility:** Jenkins offers a rich ecosystem of plugins that extend its functionality, allowing users to customize and adapt Jenkins to suit their specific requirements.

Jenkins CI/CD Pipelines:

Jenkins CI/CD pipelines define the stages and actions to be executed as part of the software delivery process. A typical pipeline includes stages for building, testing, and deploying applications, with each stage consisting of one or more tasks or jobs.

Access Jenkins: <http://localhost:8080>

Create a Pipeline Project in Jenkins



The screenshot shows the Jenkins 'Create New Item' dialog box. At the top, there is a text input field labeled 'Enter an item name' containing the text 'Review Radar'. Below this, there is a note '» Required field'. The dialog lists several project types with their descriptions and icons:

- Freestyle project**: Classic, general-purpose job type that checks out from up to one SCM, executes build steps serially, followed by post-build steps like archiving artifacts and sending email notifications.
- Pipeline**: Orchestrates long-running activities that can span multiple build agents. Suitable for building pipelines (formerly known as workflows) and/or organizing complex activities that do not easily fit in free-style job type.
- Multi-configuration project**: Suitable for projects that need a large number of different configurations, such as testing on multiple environments, platform-specific builds, etc.
- Folder**: Creates a container that stores nested items in it. Useful for grouping things together. Unlike view, which is just a filter, a folder creates a separate namespace, so you can have multiple things of the same name as long as they are in different folders.
- Multibranch Pipeline**: Creates a set of Pipeline projects according to detected branches in one SCM repository.
- Organization Folder**: Creates a set of multibranch project subfolders by scanning for repositories.

At the bottom of the dialog, there is a note: 'If you want to create a new item from other existing, you can use this option:' followed by an 'OK' button.

Project Configuration Details

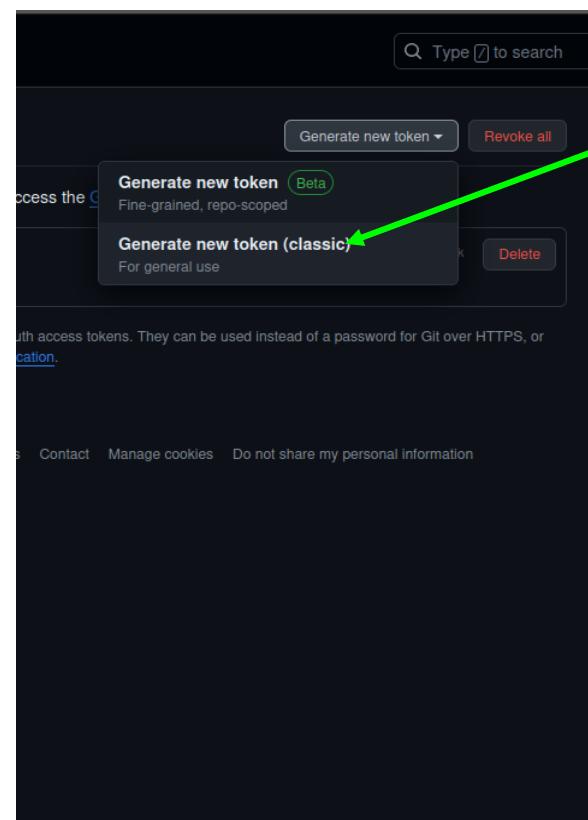
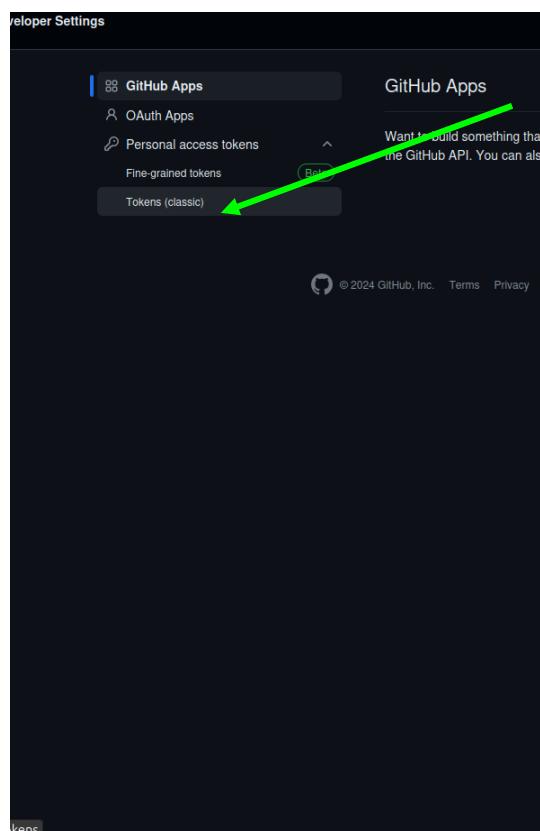
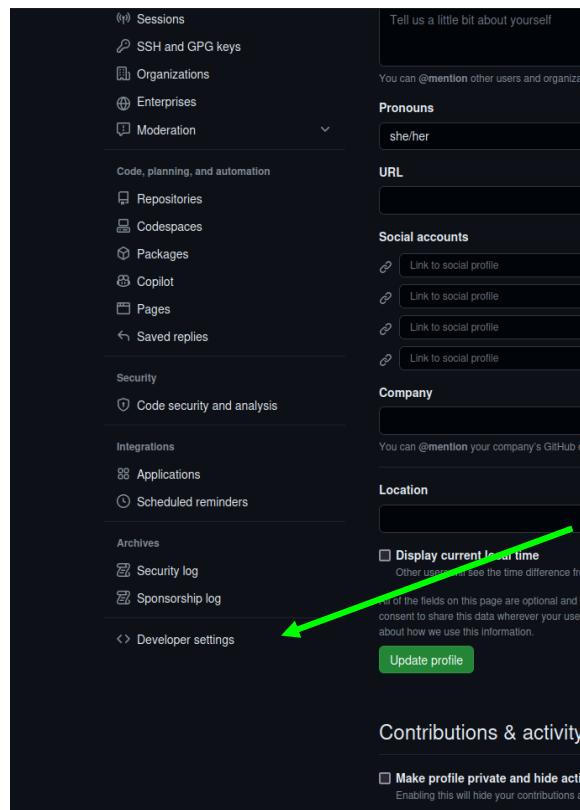
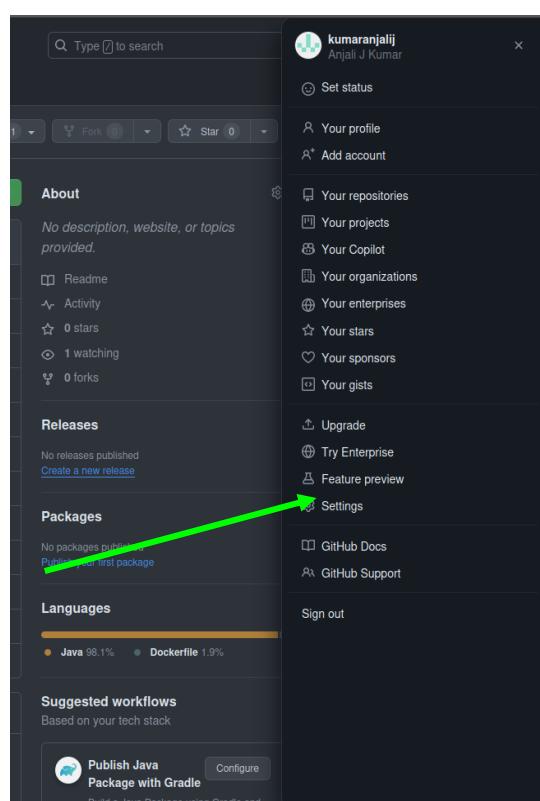
1. GitHub hook trigger for GITScm polling:

- Enabled GitHub hook trigger for GITScm polling in Jenkins to enable Git SCM polling for continuous integration.
- This configuration facilitates automatic triggering of pipeline builds upon code changes in the associated GitHub repository.

The screenshot shows the Jenkins Pipeline configuration page under the 'General' tab. In the 'Build Triggers' section, the checkbox for 'GitHub hook trigger for GITScm polling' is checked, while other options like 'Poll SCM' and 'Trigger builds remotely' are unchecked. Other tabs visible include 'Advanced Project Options' and 'Pipeline'.

2. Integration with GitHub Repository with Jenkins:

- Established seamless integration between Jenkins and project GitHub repository.
- Git is the Source Control Management tool for version control.
- Secured the integration by generating a token within GitHub with requisite permissions (repo and admin: repo hook).
- Used the token as credentials for authentication, ensuring access to the private repository.



Tokens (classic)

Note
webhook-token

What's this token for?

Expiration *
30 days The token will expire on Sun, Mar 24 2024

Select scopes
Scopes define the access for personal tokens. [Read more about OAuth scopes.](#)

<input checked="" type="checkbox"/> repo	Full control of private repositories
<input checked="" type="checkbox"/> repo:status	Access commit status
<input checked="" type="checkbox"/> repo_deployment	Access deployment status
<input checked="" type="checkbox"/> public_repo	Access public repositories
<input checked="" type="checkbox"/> repo:invite	Access repository invitations
<input checked="" type="checkbox"/> security_events	Read and write security events
<input type="checkbox"/> workflow	Update GitHub Action workflows
<input type="checkbox"/> write:packages	Upload packages to GitHub Package Registry
<input type="checkbox"/> read:packages	Download packages from GitHub Package Registry
<input type="checkbox"/> delete:packages	Delete packages from GitHub Package Registry
<input type="checkbox"/> admin:org	Full control of orgs and teams, read and write org projects
<input type="checkbox"/> write:org	Read and write org and team membership, read and write org projects
<input type="checkbox"/> read:org	Read org and team membership, read org projects
<input type="checkbox"/> manage_runners:org	Manage org runners and runner groups
<input type="checkbox"/> admin:public_key	Full control of user public keys
<input type="checkbox"/> write:public_key	Write user public keys
<input type="checkbox"/> read:public_key	Read user public keys
<input checked="" type="checkbox"/> admin:repo_hook	Full control of repository hooks
<input checked="" type="checkbox"/> write:repo_hook	Write repository hooks
<input checked="" type="checkbox"/> read:repo_hook	Read repository hooks
<input type="checkbox"/> admin:org_hook	Full control of organization hooks

Generate token Cancel

© 2024 GitHub, Inc. Terms Privacy Security Status Docs Contact Manage cookies Do not share my personal information

Personal access tokens (classic)

Generate new token Revoke all

Tokens you have generated that can be used to access the [GitHub API](#).

Make sure to copy your personal access token now. You won't be able to see it again!

ghp_OAAPrRj1qrfw600sgFRhv1MrwHPeGc01YBGA 

Delete

Copy and save the token for future use.

3. Pipeline script from SCM:

- Selected the "Pipeline script from SCM" option in Jenkins.
- Specified Git as the SCM tool.
- Provided the GitHub repository URL in Jenkins configuration.
- Configured Jenkins to authenticate with the GitHub repository using the created webhook-token.
- Defined the script path by naming the pipeline script file (jenkinsfile), ensuring proper execution of the pipeline upon triggering.

The screenshot shows the Jenkins Pipeline configuration page. The left sidebar has tabs for General, Advanced Project Options, and Pipeline, with Pipeline selected. The main area is titled 'Pipeline' under 'Definition' and shows 'Pipeline script from SCM'. The 'SCM' section is set to 'Git'. Under 'Repositories', the 'Repository URL' is set to 'https://github.com/kumaranjalj/ReviewRadar.git' and the 'Credentials' are 'kumaranjalj/******** (Review Radar webhook token)'. There is an 'Advanced' dropdown and a 'Add Repository' button. Under 'Branches to build', the 'Branch Specifier' is set to '*/*main'. In the bottom section, 'Repository browser' is set to '(Auto)', 'Additional Behaviours' has an 'Add' button, 'Script Path' is set to 'Jenkinsfile', and 'Lightweight checkout' is checked. At the bottom are 'Save' and 'Apply' buttons.

Git Plugins

The screenshot shows the Jenkins 'Plugins' page with a search bar containing 'git'. The results are listed below, ordered by name:

- Git 5.2.1**: This plugin integrates Git with Jenkins. [Report an issue with this plugin](#)
- Git client plugin 4.7.0**: Utility plugin for Git support in Jenkins. [Report an issue with this plugin](#)
- GitHub 1.38.0**: This plugin integrates GitHub to Jenkins. [Report an issue with this plugin](#)
- GitHub API Plugin 1.318-461.v7a_c09c9fa_d63**: This plugin provides GitHub API for other plugins. [Report an issue with this plugin](#)
- Github Branch Source 1785.v99802b_69816c**: Multibranch projects and organization folders from GitHub. Maintained by CloudBees, Inc. [Report an issue with this plugin](#)
- Pipeline: GitHub Groovy Libraries 42.v0739460cda_c4**: Allows Pipeline Groovy libraries to be loaded on the fly from GitHub. [Report an issue with this plugin](#)

Pipeline Script for Continuous Integration

This Jenkins pipeline automates the process of checking out code, building and testing the backend, building the frontend, creating Docker images, pushing these images to Docker Hub, and finally deploying the application using Ansible. The stages of the pipeline are as follows:

1. GIT CHECKOUT:

- This stage checks out the code from the specified GitHub repository. It uses the **git** command with the branch **main** and credentials identified by **ReviewRadar-webhook**. This ensures that the latest code from the repository is available for subsequent stages.

2. BACKEND BUILD AND TEST:

- In this stage, the backend application is built and tested. It activates a Python virtual environment and installs necessary dependencies listed in the **requirements.txt** file. The Natural Language Toolkit (NLTK) **vader_lexicon** and spaCy's **en_core_web_sm** model are also downloaded. After the dependencies are installed, the backend application is started using **nohup** to run in the background, with logs directed to **app.log**.

3. FRONTEND BUILD:

- This stage handles the building of the frontend application. It changes to the frontend project directory (**./frontend/reviewradar/**) and runs **npm install** to install

dependencies. Subsequently, **npm run build** is executed to create a production build of the frontend application.

```

1  pipeline {
2      agent any
3
4      environment{
5          DOCKER_IMAGE_NAME_FRONTEND = 'anjalijkumar/review-radar-frontend'
6          DOCKER_IMAGE_NAME_BACKEND = 'anjalijkumar/review-radar-backend'
7          GITHUB_REPO_URL = 'https://github.com/kumaranjalij/ReviewRadar.git'
8          DOCKERHUB_CREDENTIALS = credentials('DockerHubCred')
9      }
10
11      stages {
12          stage('GIT CHECKOUT') {
13              steps {
14                  script {
15                      // Checkout the code from the GitHub repo
16                      git branch: 'main', credentialsId: 'ReviewRadar-webhook', url: "${GITHUB_REPO_URL}"
17                  }
18              }
19          }
20
21
22          stage('BACKEND BUILD AND TEST') {
23              steps {
24                  // Activate the virtual environment
25                  withEnv(['PATH+VENV=/var/lib/jenkins/workspace/ReviewRadar/rrenv/bin']) {
26                      // Install packages from requirements.txt within the virtual environment
27                      sh 'python3 -m nltk.download vader_lexicon'
28                      sh 'python3 -m spacy download en_core_web_sm'
29
30                      sh 'pip install -r ./backend/Review_Radar/requirements.txt'
31
32                      // Run your Python script within the virtual environment
33                      sh 'nohup python3 ./backend/Review_Radar/app.py > app.log 2>&1 &'
34                  }
35              }
36          }
37      }

```

```
57
58     stage('FRONTEND BUILD') {
59         steps {
60             // Change directory to your specific directory
61             dir("./frontend/reviewradar") {
62                 // Run npm install and npm run build commands
63                 sh 'npm install'
64                 sh 'npm run build'
65             }
66         }
67     }
68
69
70     stage('BUILD DOCKER IMAGES') {
71         steps {
72             script {
73                 // Build and tag backend Docker image
74                 docker.build("${DOCKER_IMAGE_NAME_BACKEND}", './backend/Review_Radar')
75                 sh 'docker build -t "${DOCKER_IMAGE_NAME_BACKEND}" ./backend/Review_Radar'
76
77                 // Build and tag frontend Docker image
78                 docker.build("${DOCKER_IMAGE_NAME_FRONTEND}", './frontend/reviewradar')
79                 sh 'docker build -t "${DOCKER_IMAGE_NAME_FRONTEND}" ./frontend/reviewradar'
80             }
81         }
82     }
83
84     stage('PUSH DOCKER IMAGE') {
85         steps {
86             script{
87                 sh 'echo "Logging in to Docker Hub"'
88                 sh """
89                 echo \$DOCKERHUB_CREDENTIALS_PSW | docker login -u \$DOCKERHUB_CREDENTIALS_USR --password-stdin
90                 """
91
92                 sh 'echo "pushing frontend"'
93                 sh 'docker push anjalijkumar/review-radar-frontend:latest'
94
95                 sh 'echo "pushing backend"'
96                 sh 'docker push anjalijkumar/review-radar-backend:latest'
97             }
98         }
99     }
100
101    stage('RUN ANSIBLE PLAYBOOK') {
102        steps {
103            script {
104                ansiblePlaybook(
105                    playbook: './ansible-deploy/deploy.yml',
106                    inventory: './ansible-deploy/inventory'
107                )
108            }
109        }
110    }
111 }
```

4. BUILD DOCKER IMAGES:

- In this stage, Docker images for both the backend and frontend applications are built. The backend image is created using the Dockerfile located in `./backend/Review_Radar/`, and the frontend image is built from the Dockerfile in `./frontend/reviewradar/`. Each image is tagged appropriately using the environment variables **DOCKER_IMAGE_NAME_BACKEND** and **DOCKER_IMAGE_NAME_FRONTEND**.

5. PUSH DOCKER IMAGE:

- This stage pushes the built Docker images to Docker Hub. First, it logs into Docker Hub using credentials stored in the environment variables **DOCKERHUB_CREDENTIALS_USR** and **DOCKERHUB_CREDENTIALS_PSW**. Then, it pushes the frontend and backend images to their respective repositories on Docker Hub.

6. RUN ANSIBLE PLAYBOOK:

- The final stage runs an Ansible playbook to deploy the application onto a Kubernetes cluster. The **ansiblePlaybook** function is used to execute the playbook defined in `./ansible-deploy/deploy.yml` with the inventory file located at `./ansible-deploy/inventory`. This stage ensures that the latest versions of the application are deployed and running in the specified environment.

Ngrok

Ngrok provides a secure way to share and test web applications by creating temporary, publicly accessible URLs that can be accessed from anywhere.

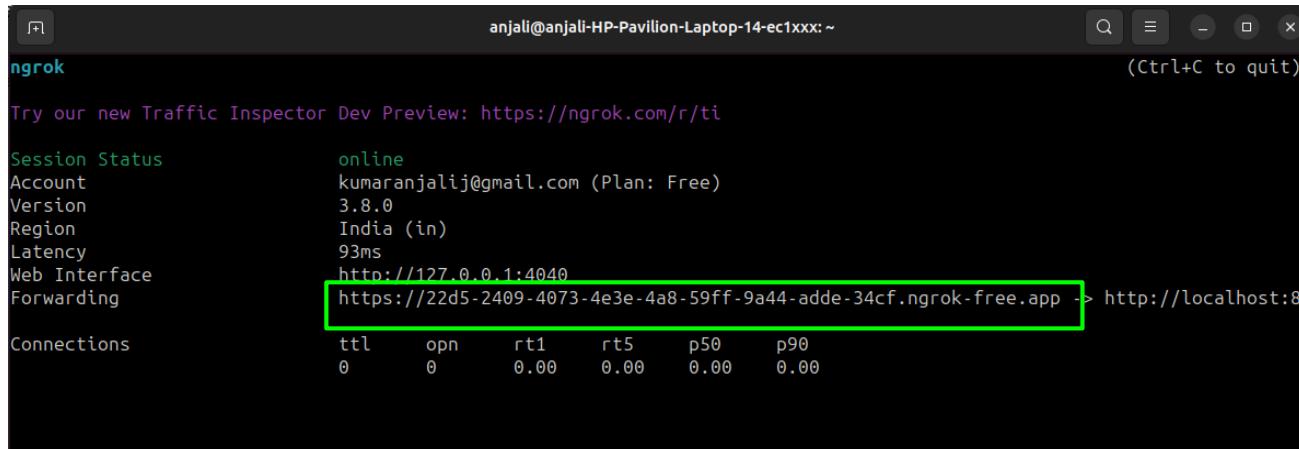
Installation

- Sign up in <https://ngrok.com/>
- Download ngrok from:
<https://bin.equinox.io/c/4VmDzA7iaHb/ngrok-stable-linux-amd64.tgz>
- Then extract ngrok from the terminal:
`$sudo tar xvzf ~/Downloads/ngrok-stable-linux-amd64.tgz -C /usr/local/bin`
- Copy Authtoken from:
<https://dashboard.ngrok.com/get-started/your-authtoken>
- Add Authtoken:
`$ngrok authtoken <token>`

How to get public IP address

Execute the following command and copy the public ip address for your local host.

```
$ngrok http 8080
```



```

ngrok
Try our new Traffic Inspector Dev Preview: https://ngrok.com/r/ti

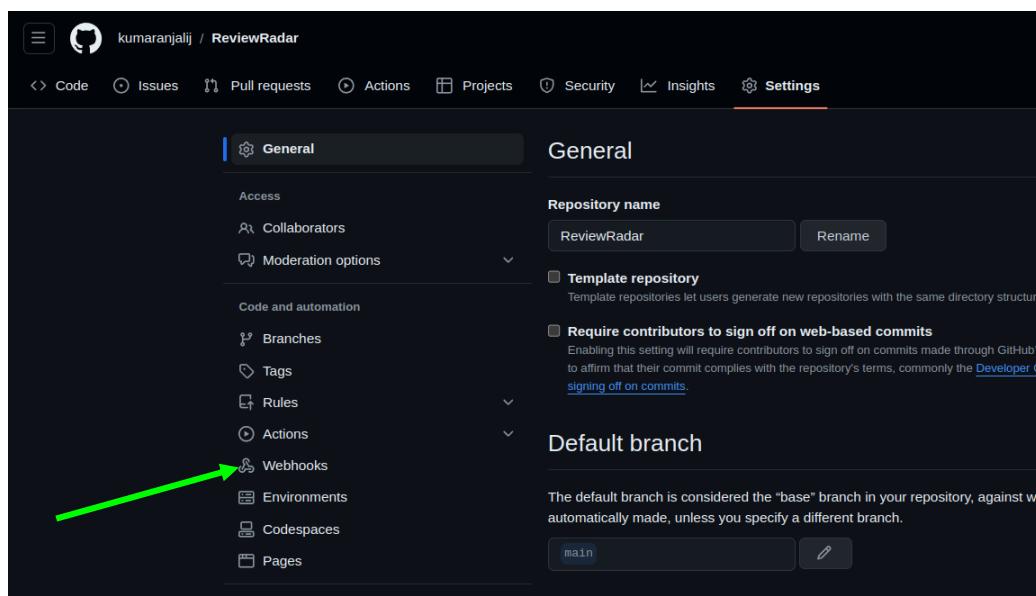
Session Status          online
Account                 kumaranjalij@gmail.com (Plan: Free)
Version                3.8.0
Region                 India (in)
Latency               93ms
Web Interface          http://127.0.0.1:4040
Forwarding             https://22d5-2409-4073-4e3e-4a8-59ff-9a44-adde-34cf.ngrok-free.app -> http://localhost:8
Connections            ttl     opn     rtt1    rtt5     p50     p90
                        0       0      0.00   0.00   0.00   0.00

```

GitHub Webhook

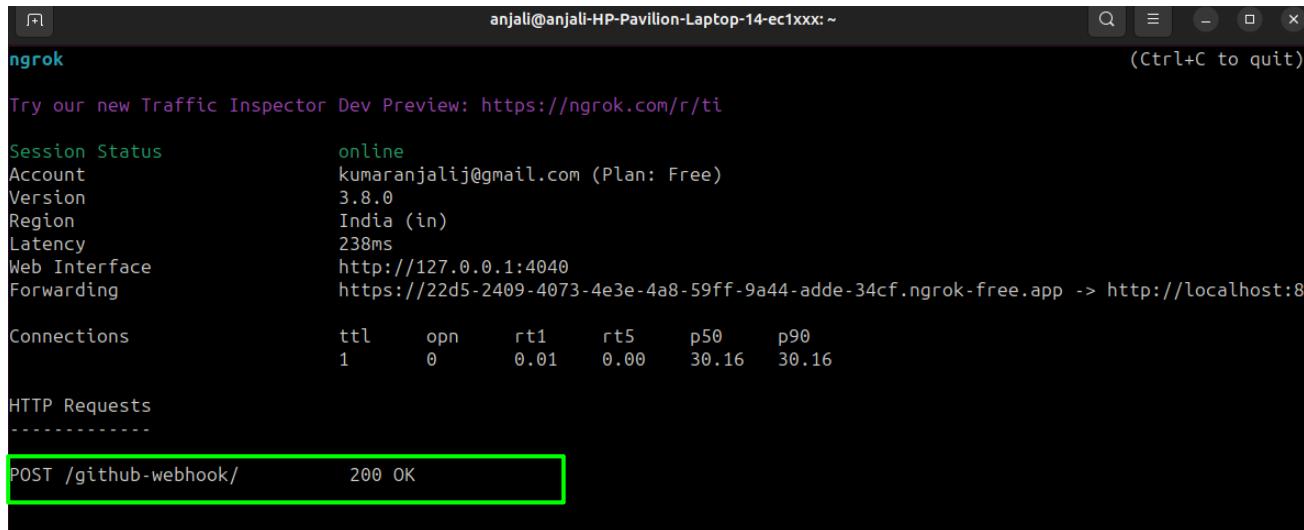
Add a webhook in the GitHub repository with ngrok's public URL .The webhook-token created earlier is used as the secret key.

- Navigate to your GitHub repository.
- Go to the repository's settings.
- Select "Webhooks" or "Webhooks & services" from the menu.
- Click on "Add webhook".
- In the "Payload URL" field, paste the ngrok public URL.
- Set the "Content type" to **application/json**.
- In the "Secret" field, paste the webhook-token.
- Choose the events you want the webhook to trigger on. Typically, this includes events like push, pull request, etc.
- Optionally, you can set other configurations such as SSL verification and webhook status.



The screenshots illustrate the process of setting up a webhook in GitHub:

- Step 1: Initiate Setup**
The first screenshot shows the GitHub settings page for the repository "kumaranjalij / ReviewRadar". The "Webhooks" section is selected. A green arrow points to the "Add webhook" button at the top right of the "Webhooks" section.
- Step 2: Configure Webhook**
The second screenshot shows the "Webhooks / Add webhook" configuration form. It includes fields for "Payload URL" (containing the URL "https://73-4e3e-4a8-59ff-9a44-adde-34cf.ngrok-free.app/github-webhook/"), "Content type" (set to "application/json"), "Secret" (containing the value "ghp_uZValeYKonPKI2ozh3H4d4Oz3limyO0bqFkw"), and "SSL verification" (with "Enable SSL verification" checked). A green arrow points to the "Add webhook" button at the bottom left.
- Step 3: View Webhooks**
The third screenshot shows the "Webhooks" list page. It displays a single webhook entry with a green checkmark icon, the URL "https://22d5-2409-4073-4e3e-4a8-...", and the event type "(push)". To the right of the URL are "Edit" and "Delete" buttons. A green arrow points to the webhook entry.



```

anjali@anjali-HP-Pavilion-Laptop-14-ec1xxx: ~
ngrok
Try our new Traffic Inspector Dev Preview: https://ngrok.com/r/ti

Session Status      online
Account            kumaranjalij@gmail.com (Plan: Free)
Version             3.8.0
Region              India (in)
Latency             238ms
Web Interface      http://127.0.0.1:4040
Forwarding          https://22d5-2409-4073-4e3e-4a8-59ff-9a44-adde-34cf.ngrok-free.app -> http://localhost:8
Connections         ttl     opn     rt1     rt5     p50     p90
                    1       0     0.01   0.00   30.16   30.16

HTTP Requests
-----
POST /github-webhook/    200 OK

```

Configure ngrok in jenkins

Configure ngrok URL, GitHub API URL, and provide webhook token as credentials in Jenkins, to ensure that Jenkins is correctly set up to communicate with GitHub via ngrok, and it has the necessary authentication credentials to interact with GitHub's webhook system securely. To configure ngrok in Jenkins and manage webhook settings for GitHub integration, you would follow these steps:

1. Configure Ngrok URL in Jenkins:

- Navigate to "Manage Jenkins" > "System".
- Look for the "Jenkins URL" field and paste the ngrok public URL.

Jenkins Location

Jenkins URL	https://22d5-2409-4073-4e3e-4a8-59ff-9a44-adde-34cf.ngrok-free.app
System Admin e-mail address	Jenkins-Master <kumaranjalij@gmail.com>

2. GitHub API URL Configuration:

- On the same page, you may have an option to configure GitHub related settings.
- Find the field for "GitHub API URL" and ensure it's correctly configured.
- Usually, it should be set to <https://api.github.com>.

GitHub

GitHub Servers ?

GitHub Server ?

Name ?
github

API URL ?
https://api.github.com

Credentials ?
- none -

+ Add ▾

**3. Configure Credentials:**

- Click on "Add" to add new credentials.
- Choose the type of credentials as “Secret Text”.
- Enter the webhook token in the "Secret" field.
- Optionally, provide an ID and description for the credentials.

Jenkins Credentials Provider: Jenkins

Add Credentials

Domain
Global credentials (unrestricted)

Kind
Secret text

Scope ?
Global (Jenkins, nodes, items, all child items, etc)

Secret
.....

ID ?
ngrokSecretCred

Description ?
Secret Cred for [ngrok](#)

Cancel Add

4. Save Changes:

- Once you've configured the ngrok URL, GitHub API URL, and added the webhook token as credentials, click on "Save" or "Apply" to save your changes.

GitHub

GitHub Servers ?

≡ GitHub Server ?

Name ?

github

API URL ?

https://api.github.com

Credentials ?

Secret Cred for ngrok

+ Add ▾

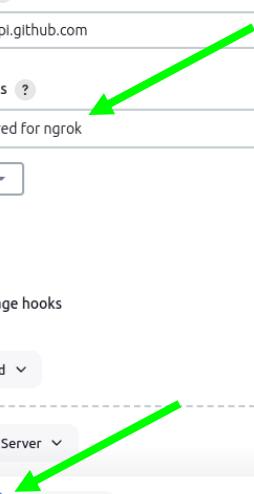
Manage hooks

Advanced ▾

Add GitHub Server ▾

Save

Apply



CONTAINERIZATION WITH DOCKER

Docker is a platform for developing, shipping, and running applications inside containers. Containerization allows applications and their dependencies to be packaged together in a consistent environment, ensuring that they run reliably across different computing environments. By containerizing both the frontend and backend components of your application, you can streamline deployment, improve scalability, and enhance portability.

Docker Plugins

The screenshot shows the Jenkins Plugins page. The navigation bar at the top includes 'Dashboard', 'Manage Jenkins', and 'Plugins'. Below the navigation, there's a search bar with the text 'Docker'. On the left, a sidebar has tabs for 'Updates' (24), 'Available plugins', 'Installed plugins' (which is selected and highlighted in grey), and 'Advanced settings'. The main content area lists several Docker-related plugins:

- Docker API Plugin** 3.3.4-86.v39b_a_5ede342c: Provides the docker-java API for other plugins. Status: Up for adoption! (Yellow box)
- Docker Commons Plugin** 439.va_3cb_0a_6a_fb_29: Provides the common shared functionality for various Docker-related plugins.
- Docker Pipeline** 572.v950f58993843: Build and use Docker containers from pipelines.
- Docker plugin** 1.6.1: Integrates Jenkins with Docker.

Creating Docker Containers:

Backend Container:

Docker containers encapsulate the backend Flask application along with its dependencies, such as libraries and configurations. This container exposes the Flask application's endpoints, allowing it to receive requests and respond accordingly.

```
1  # Use an official Python runtime as a parent image
2  FROM python:3.8-slim
3
4  # Install system-level dependencies
5  RUN apt-get update \
6      && apt-get install -y pkg-config \
7          libmariadb-dev-compat \
8          gcc \
9      && rm -rf /var/lib/apt/lists/*
10
11 # Set the working directory in the container
12 WORKDIR /app
13
14 # Copy the requirements file into the container at /app
15 COPY requirements.txt /app/
16
17 # Install any needed dependencies specified in requirements.txt
18 RUN pip install --no-cache-dir -r requirements.txt
19
20 RUN python -c "import nltk; nltk.download('vader_lexicon')"
21
22 RUN python -m spacy download en_core_web_sm
23
24 # Copy the rest of the application code into the container at /app
25 COPY . /app/
26
27 # Make port 5001 available to the world outside this container
28 EXPOSE 5001
29
30 # Define environment variable
31 ENV NAME World
32
33 # Run app.py when the container launches
34 CMD ["python3", "app.py"]
```

Frontend Container:

The frontend React application is packaged into a separate Docker container, which contains the necessary libraries, assets, and configurations. This container serves the static files and assets required to render the frontend user interface.

```

1   # Use an official Node.js runtime as a parent image
2   FROM node:14-alpine
3
4   # Set the working directory in the container
5   WORKDIR /app
6
7   # Copy package.json and package-lock.json to the container
8   COPY package*.json ./
9
10  # Install dependencies
11  RUN npm install
12  RUN npm install axios
13  RUN npm install react-router-dom
14
15  # Copy the rest of the application code to the container
16  COPY . .
17
18  # Expose port 3000
19  EXPOSE 3000
20
21  # Command to start serving the React application
22  CMD ["npm", "start"]

```

Setting Global Credentials

Dashboard > Manage Jenkins > Credentials

Credentials

T	P	Store	Domain	ID	Name
System	System	(global)	ReviewRadar	kumaranjalij/******** (ReviewRadar token)	
System	System	(global)	ReviewRadar-webhook	kumaranjalij/******** (Review Radar webhook token)	
System	System	(global)	DockerHubCred	anjaliukumar/******** (Dockerhub Credentials)	
System	System	(global)	localhost	anjali/******** (localhost Credentials)	
System	System	(global)	ngrokSecretCred	Secret Cred for ngrok	

Stores scoped to Jenkins

DEPLOYMENT WITH KUBERNETES

Kubernetes is a powerful container orchestration platform that automates the deployment, scaling, and management of containerized applications. It provides features such as service discovery, load balancing, and automated rollouts and rollbacks, which are essential for deploying applications in production environments. By using Kubernetes, you can ensure high availability, scalability, and reliability for your containerized application.

Kubernetes Plugins

The screenshot shows the Jenkins Plugins page. The left sidebar has tabs for 'Updates' (24), 'Available plugins', 'Installed plugins' (selected), and 'Advanced settings'. The main area has a search bar with 'Kubernetes' typed in. Below it, a table lists several Kubernetes-related plugins:

Name	Version	Description
Kubernetes :: Pipeline :: DevOps Steps	1.6	Report an issue with this plugin
Kubernetes CLI Plugin	1.12.1	Configure kubectl for Kubernetes Report an issue with this plugin
Kubernetes Client API Plugin	6.10.0-240.v57880ce8b_0b_2	Kubernetes Client API plugin for use by other Jenkins plugins. Report an issue with this plugin
Kubernetes Credentials Plugin	0.11	Common classes for Kubernetes credentials Report an issue with this plugin
Kubernetes Credentials Provider	1.262.v2670ef7ea_0c5	Provides a read only credentials store backed by Kubernetes. Report an issue with this plugin
Kubernetes plugin	4219.v40ff98cfb_d6f	This plugin integrates Jenkins with Kubernetes Report an issue with this plugin

Kubernetes Deployment Configuration:

Backend deploy

This Kubernetes Deployment configuration defines a deployment named "review-radar-backend" with one replica, meaning it ensures that one instance of the backend application is running. Within this deployment, a container named "review-radar-backend" is specified. This container is based on the Docker image "anjalijkumar/review-radar-backend:latest", which contains the necessary code and dependencies for the backend application. The deployment exposes port 5001, allowing other components within the Kubernetes cluster to communicate with the backend application. This configuration ensures the reliable deployment and availability of the backend application within the Kubernetes environment, facilitating seamless communication and interaction with other services or components.

```

1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: review-radar-backend
5  spec:
6    replicas: 1
7    selector:
8      matchLabels:
9        app: review-radar-backend
10   template:
11     metadata:
12       labels:
13         app: review-radar-backend
14   spec:
15     containers:
16       - name: review-radar-backend
17         image: anjalijkumar/review-radar-backend:latest
18         ports:
19           - containerPort: 5001

```

Frontend Deploy

```

1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: review-radar-frontend
5  spec:
6    replicas: 1
7    selector:
8      matchLabels:
9        app: review-radar-frontend
10   template:
11     metadata:
12       labels:
13         app: review-radar-frontend
14   spec:
15     containers:
16       - name: review-radar-frontend
17         image: anjalijkumar/review-radar-frontend:latest
18         ports:
19           - containerPort: 3000

```

This Kubernetes Deployment configuration defines a deployment named "review-radar-frontend" with one replica, ensuring the availability of one instance of the frontend application. Within this deployment, a container named "review-radar-frontend" is specified, which is based on the Docker image "anjalijkumar/review-radar-frontend:latest". This Docker image contains the necessary code and dependencies for the frontend application. The deployment exposes port 3000, allowing other components within the Kubernetes cluster to communicate with the frontend application. Overall,

this configuration facilitates the reliable deployment and accessibility of the frontend application within the Kubernetes environment, enabling seamless interaction with users and other services.

Kubernetes Service Configuration:

Backend service config

```
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: review-radar-backend-service
5  spec:
6    selector:
7      app: review-radar-backend
8    ports:
9      - protocol: TCP
10     port: 5001
11     targetPort: 5001
12   type: NodePort
```

This Kubernetes Service configuration defines a service named "review-radar-backend-service" that acts as a bridge between clients and the backend application deployed within the Kubernetes cluster. The service ensures that external or internal components can communicate with the backend application seamlessly.

The service is configured to route traffic to pods with the label "app: review-radar-backend", meaning it directs requests to instances of the backend application. This selector ensures that the service forwards traffic specifically to pods running the backend application.

The service exposes port 5001, which is the port on which the backend application is listening inside the pods. When requests are received on port 5001, the service forwards them to the corresponding pods.

Additionally, the service specifies that it uses the TCP protocol and maps port 5001 of the service to port 5001 of the backend application pods using the "targetPort" attribute.

This configuration ensures that the backend application is accessible within the Kubernetes cluster through the "review-radar-backend-service", allowing other components or services to interact with it reliably.

Frontend service config

```
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: review-radar-frontend-service
5  spec:
6    selector:
7      app: review-radar-frontend
8    ports:
9      - protocol: TCP
10     port: 3000
11     targetPort: 3000
12   type: NodePort
```

This Kubernetes Service configuration defines a service named "review-radar-frontend-service" responsible for routing traffic to the frontend application deployed within the Kubernetes cluster. The service acts as an entry point, enabling external or internal clients to communicate with the frontend application seamlessly.

The service is configured to select pods with the label "app: review-radar-frontend", ensuring that it directs incoming traffic specifically to instances of the frontend application.

It exposes port 3000, which is the port on which the frontend application is listening inside the pods. Incoming requests received on port 3000 are forwarded to the corresponding pods running the frontend application.

This service uses the TCP protocol to communicate with the backend application, and it maps port 3000 of the service to port 3000 of the frontend application pods using the "targetPort" attribute.

The configuration facilitates reliable access to the frontend application within the Kubernetes cluster through the "review-radar-frontend-service", allowing other components or services to interact with it as needed.

Ingress File

```

1      apiVersion: networking.k8s.io/v1
2      kind: Ingress
3      metadata:
4          name: reviewradar-ingress
5          annotations:
6              ingress.kubernetes.io/rewrite-target: /
7              kubernetes.io/ingress.class: "nginx"
8      spec:
9          ingressClassName: nginx
10         rules:
11             - host: localhost
12                 http:
13                     paths:
14                         - path: /
15                             pathType: Prefix
16                             backend:
17                                 service:
18                                     name: review-radar-backend-service
19                                     port:
20                                         number: 5001
21                         - path: /
22                             pathType: Prefix
23                             backend:
24                                 service:
25                                     name: review-radar-frontend-service
26                                     port:
27                                         number: 3000

```

This Kubernetes Ingress configuration facilitates external access to services within the cluster through defined routing rules:

Annotations: These provide additional configuration options for the Ingress controller. For instance:

- "ingress.kubernetes.io/rewrite-target": Rewrites incoming request paths to "/".
- "kubernetes.io/ingress.class": Specifies the Ingress controller class as "nginx".

Spec

- **IngressClassName:** Specifies the Ingress controller class as "nginx".
- **Rules:** These define routing based on hostname and path:
 - **Host:** Identifies the hostname to which the rules apply, in this case, "localhost".
 - **HTTP Paths:** Define routing based on path prefix:
 - **Path:** Specifies the path prefix.

- **PathType:** Indicates the type of matching used for the path, here, "Prefix".
- **Backend:** Specifies the backend service to which requests are forwarded:
 - **Service Name:** Name of the backend service.
 - **Port Number:** Port of the backend service

This setup ensures that incoming HTTP requests are directed to the appropriate backend services based on defined rules, enabling efficient external access to services within the Kubernetes cluster.

CONFIGURATION MANAGEMENT WITH ANSIBLE

Ansible is an open-source automation tool that enables configuration management, application deployment, and orchestration. It allows you to automate tasks such as software installation, configuration changes, and system updates across multiple servers or environments. Ansible uses simple YAML-based playbooks to define automation tasks, making it easy to understand and use.

Benefits of Ansible:

- **Consistency:** Ansible ensures consistency across environments by automating configuration tasks and enforcing desired states on target systems.
- **Scalability:** Ansible scales easily to manage configurations for large infrastructure deployments, enabling efficient management of multiple servers.
- **Reduced Manual Overhead:** By automating repetitive tasks, Ansible reduces manual overhead and frees up valuable time for IT teams to focus on more strategic initiatives.
- **Idempotent Operations:** Ansible playbooks are idempotent, meaning they can be run multiple times without causing unexpected changes, ensuring predictable and reliable configuration management.

Ansible Plugins

The screenshot shows the Jenkins management interface for plugins. The left sidebar has tabs for 'Updates' (24), 'Available plugins', 'Installed plugins' (selected), and 'Advanced settings'. A search bar at the top contains the text 'Ansible'. Below the search bar, a table lists two plugins:

Name
Ansible plugin 307.va_1f3ef06575a_
Ansible Tower Plugin 0.16.0

 Each row includes a link to invoke the plugin and another to report an issue.

Ansible Playbook:

Ansible playbooks are YAML files that define a set of tasks to be executed on remote systems. Each task specifies the desired state of a system, such as installing packages, configuring services, or updating configuration files.

```

1   ---
2   - name: Deploying Review Radar on the kubernetes cluster
3     hosts: localhost
4     remote_user: anjali
5     become: false
6     tasks:
7       - name: Deleting older application files
8         ansible.builtin.shell: |
9           kubectl delete -f ../k8s/deploy/backend-deploy.yml --ignore-not-found=true
10          kubectl delete -f ../k8s/deploy/frontend-deploy.yml --ignore-not-found=true
11          kubectl delete -f ../k8s/services/backend-service.yml --ignore-not-found=true
12          kubectl delete -f ../k8s/services/frontend-service.yml --ignore-not-found=true
13          kubectl delete -f ../k8s/ingress/ingress.yml --ignore-not-found=true
14
15       - name: Deploying application to kubernetes cluster
16         ansible.builtin.shell: |
17           kubectl apply -f ../k8s/deploy/backend-deploy.yml
18           kubectl apply -f ../k8s/deploy/frontend-deploy.yml
19           kubectl apply -f ../k8s/services/backend-service.yml
20           kubectl apply -f ../k8s/services/frontend-service.yml
21           kubectl apply -f ../k8s/ingress/ingress.yml

```

This is an Ansible playbook named "Deploying Review Radar on the Kubernetes cluster". Ansible is being used to automate the deployment process of the Review Radar application onto a Kubernetes cluster.

name: Describes the name of the playbook.

hosts: Specifies the target hosts or machines where the tasks will be executed. Here, it's set to "localhost", indicating that the tasks will be executed on the local machine.

remote_user: Specifies the user to use for SSH connections. Here, it's set to "anjali".

become: Indicates whether to escalate privileges (become root) to execute tasks. Here, it's set to "false", meaning privileges won't be escalated.

tasks: This section contains the list of tasks to be executed:

Deleting older application files: This task uses the kubectl delete command to delete older application files from the Kubernetes cluster. It includes deleting deployment files (backend-deploy.yml and frontend-deploy.yml), service files (backend-service.yml and frontend-service.yml), and the ingress file (ingress.yml). The --ignore-not-found=true option ensures that the command does not fail if the resources are not found.

Deploying application to Kubernetes cluster: This task uses the kubectl apply command to apply new application files to the Kubernetes cluster. It includes applying deployment files for backend and

frontend (backend-deploy.yml and frontend-deploy.yml), service files for backend and frontend (backend-service.yml and frontend-service.yml), and the ingress file (ingress.yml).

Overall, this playbook automates the process of deploying the Review Radar application to a Kubernetes cluster by deleting old resources and applying new ones using kubectl commands.

Inventory file

In Ansible, an inventory file is a crucial component that defines the hosts or groups of hosts on which Ansible commands, modules, and playbooks will operate. The inventory file can be a simple text file listing IP addresses or hostnames, organized into groups for easier management. It allows users to specify variables and configuration settings for each host or group, enabling efficient and targeted automation.

```
1     ansible_host=localhost ansible_user=anjali ansible_ssh_pass=12345
2     ansible_ssh_common_args = '-o StrictHostKeyChecking=no'
```

Here's a breakdown of each part:

ansible_host=localhost: Specifies the address of the host to connect to. In this case, localhost indicates the local machine.

ansible_user=anjali: Indicates the username to use when connecting to the host. Here, the user is anjali.

ansible_ssh_pass=12345: Provides the SSH password for the user anjali. In this example, the password is 12345

ansible_ssh_common_args='-o StrictHostKeyChecking=no': Sets additional SSH arguments. The **-o StrictHostKeyChecking=no** option disables SSH's host key checking, preventing SSH from prompting for confirmation when connecting to a new host for the first time. This is useful for automation but should be used with caution as it can expose the system to man-in-the-middle attacks.

This configuration provides Ansible with the necessary details to connect to localhost using the specified user credentials and SSH options.

Jenkins Pipeline Output

The screenshot shows the Jenkins Pipeline Output for the ReviewRadar project. The sidebar on the left provides navigation and management options. The main area is titled "Stage View" and displays a grid of execution times for various stages across two recent builds. The grid columns represent different stages: Declarative: Checkout SCM, GIT CHECKOUT, BACKEND BUILD AND TEST, FRONTEND BUILD, BUILD DOCKER IMAGES, PUSH DOCKER IMAGE, and RUN ANSIBLE PLAYBOOK. The rows show build details like commit count and execution time.

This Jenkins pipeline provides an automated, consistent, and repeatable process for building, testing, containerizing, and deploying the Review Radar application. By breaking down the process into distinct stages, it ensures each step is executed in sequence, with clear checkpoints and validations along the way. This approach enhances efficiency, reduces the potential for human error, and facilitates continuous integration and continuous deployment (CI/CD) practices.

Kubernetes Dashboard

The screenshot shows the Kubernetes Dashboard Workloads page. The sidebar on the left lists various Kubernetes resources. The main area displays the status of workloads, including Deployments, Pods, and Replica Sets, represented by green circles indicating they are running 2 instances. Below this, detailed tables provide information about Deployments and Pods, including their names, images, labels, and resource usage.

Name	Images	Labels	Pods	Created
review-radar-backend	anjaliikumar/review-radar-backend:latest	-	1 / 1	a minute ago
review-radar-frontend	anjaliikumar/review-radar-frontend:latest	-	1 / 1	a minute ago

Name	Images	Labels	Node	Status	Restarts	CPU Usage (cores)	Memory Usage (bytes)	Created
review-radar-backend-5d645899bf-smzn2	anjaliikumar/review-radar-backend:latest	app: review-radar-backend pod-template-hash: 5d645899bf	minikube	Running	0	336.00m	748.25Mi	a minute ago
review-radar-frontend								

CONTINUOUS MONITORING WITH THE ELK STACK

The ELK Stack is a popular open-source log management platform that combines Elasticsearch, Logstash, and Kibana. It provides a comprehensive solution for collecting, indexing, searching, analyzing, and visualizing log data in real-time. The ELK Stack is widely used for monitoring applications, infrastructure, and system logs, enabling organizations to gain insights into their systems' health, performance, and security.

Components of the ELK Stack:

- **Elasticsearch:** Elasticsearch is a distributed search and analytics engine that stores and indexes log data efficiently. It enables fast and scalable searches across large volumes of log data, supporting real-time analysis and visualization.
- **Logstash:** Logstash is a data processing pipeline that ingests, processes, and enriches log data before indexing it into Elasticsearch. It provides input plugins to collect logs from various sources, filter plugins to parse and transform log data, and output plugins to send processed data to Elasticsearch.
- **Kibana:** Kibana is a web-based visualization and dashboarding tool that allows users to explore, analyze, and visualize log data stored in Elasticsearch. It provides interactive dashboards, charts, and graphs for monitoring and troubleshooting purposes, making it easy to identify trends, anomalies, and patterns in log data.

Continuous Monitoring Workflow:

- **Data Collection:** Log data is collected from various sources such as servers, applications, and network devices using Logstash input plugins. These plugins support protocols like syslog, filebeat, and beats for collecting logs in different formats.
- **Data Processing:** Logstash processes the collected log data by parsing, filtering, and enriching it as required. This includes extracting relevant fields, converting timestamps, and applying transformations to normalize the data.
- **Indexing:** Processed log data is indexed into Elasticsearch, where it is stored and made available for search and analysis. Elasticsearch indexes the data in a distributed and scalable manner, ensuring fast and efficient retrieval.
- **Visualization and Analysis:** Kibana is used to create visualizations and dashboards that provide insights into log data. Users can create custom dashboards with charts, graphs, and tables to monitor system performance, track error rates, and detect security incidents in real-time.

Key Features of the ELK Stack:

- **Real-time Monitoring:** The ELK Stack supports real-time monitoring of log data, allowing organizations to detect and respond to issues as they occur.
- **Scalability:** Elasticsearch and the ELK Stack can scale horizontally to handle large volumes of log data, ensuring performance and reliability even in high-traffic environments.

- **Alerting:** Kibana provides alerting capabilities to notify users of important events or anomalies in log data. Alerts can be configured based on predefined conditions and thresholds.
 - **Search and Analysis:** Elasticsearch offers powerful search and analytics capabilities, including full-text search, aggregations, and filtering, enabling users to analyze log data in-depth and extract valuable insights.

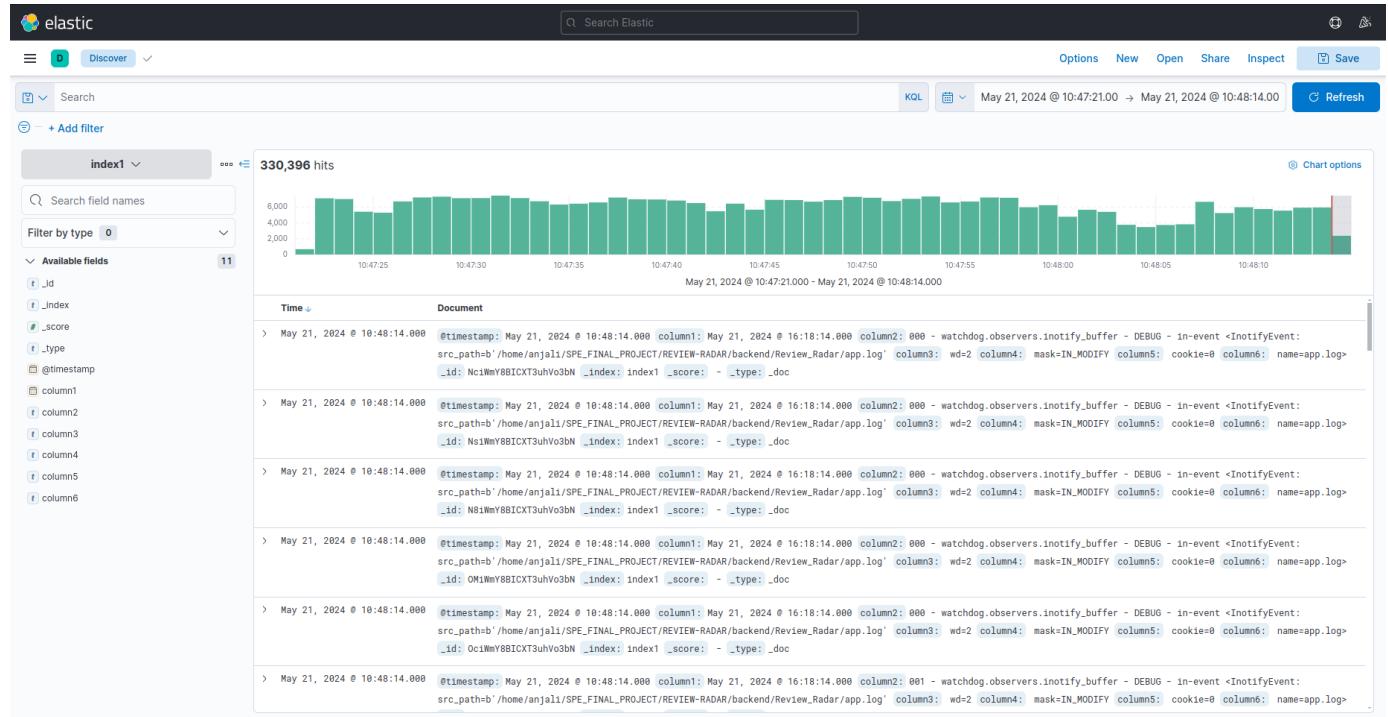
Deployment Considerations:

- **Infrastructure Requirements:** Deploying the ELK Stack requires sufficient hardware resources, especially for Elasticsearch nodes, to ensure optimal performance and scalability.
 - **Security:** Proper security measures should be implemented to protect the ELK Stack from unauthorized access and data breaches. This includes securing network communication, enabling authentication and authorization mechanisms, and encrypting sensitive data.
 - **Maintenance and Monitoring:** Ongoing maintenance and monitoring are essential for the smooth operation of the ELK Stack. Regular updates, backups, and performance monitoring help ensure the reliability and availability of the monitoring platform.

Log File

Kibana Dashboard

Access Kibana Dashboard: <http://localhost:5601>



CONCLUSION

In conclusion, the implementation of various technologies and methodologies such as topic modeling, sentiment analysis, Docker containerization, Ansible configuration management, Jenkins continuous integration, and the ELK Stack for continuous monitoring has significantly enhanced our ability to analyze customer sentiments, streamline development workflows, and ensure the reliability and performance of our applications.

Through topic modeling and sentiment analysis, we have gained valuable insights into customer feedback, enabling us to identify key areas for improvement and prioritize actionable tasks. By leveraging Docker containerization, we have achieved consistency and portability in our deployment process, allowing us to deploy applications efficiently across different environments.

Ansible configuration management has enabled us to automate repetitive tasks and enforce desired states across our infrastructure, improving consistency and reducing manual overhead. Jenkins continuous integration has facilitated the automation of build, test, and deployment processes, enabling us to deliver high-quality software faster and more reliably.

Finally, the ELK Stack for continuous monitoring has provided us with real-time visibility into our system's health, performance, and security posture. By collecting, analyzing, and visualizing log data, we can proactively detect and respond to issues, ensuring the reliability and availability of our applications.

In summary, the adoption of these technologies and methodologies has empowered us to better understand customer sentiments, accelerate development cycles, and maintain the stability and security of our applications in an ever-evolving technological landscape. As we continue to innovate and improve, these tools and practices will remain invaluable assets in our pursuit of excellence.